



KNOWLEDGE BASE COMPANION GUIDE

Your Second Brain, On Your Machine

A working library of compounding knowledge bases, with
Claude as your librarian.

*Based on Andrej Karpathy's LLM knowledge base pattern.
Adapted for CoWork OS so you can run it locally with no setup.*



THE IDEA

The Karpathy pattern

Andrej Karpathy described a simple pattern for working with LLMs over your own material: keep your raw sources untouched, let the model build a wiki from them, and answer every question from that wiki with full source provenance.

Most people throw notes, articles, and PDFs into Notion or Obsidian and hope an AI can stitch sense out of the pile. It can't. Without structure, the model invents context, loses sources, and gives you confident answers it can't back up.

Karpathy's pattern fixes that by separating three jobs into three folders:

RAW/

Your sources, word-for-word. Articles, papers, transcripts, podcasts. Never summarised. Never edited after ingest. This is the truth layer.

Wiki/

Claude reads RAW and writes the wiki — cross-linked articles, indexed, every claim tied to a source. This is the thinking layer.

Outputs/

Every question you ask becomes a report filed here, with citations back to the wiki. Strong outputs get promoted into the wiki. This is the working layer.

Why the separation matters

RAW gives you provenance. If the wiki ever drifts, you can trace any claim back to the actual words of the source. Wiki gives you synthesis — Claude reads dozens of sources, finds the through-lines, and writes them up in your voice. Outputs gives you compounding. Every question you ask leaves a permanent trail of reasoning the system can build on later.

Karpathy posted the original pattern here: x.com/karpathy/status/2039805659525644595

The thing most people get wrong about AI knowledge management: *they ask the AI to summarise on ingest. That's the moment provenance breaks. Once your raw sources are summaries of summaries, you can't trace anything back. Verbatim RAW is the rule that makes everything else trustworthy.*



THE SYSTEM

Claude as your librarian

CoWork OS ships a KNOWLEDGE/ folder at the top of your workspace. Inside, you can run as many parallel knowledge bases as you want — one per topic — and Claude maintains every one of them as you go.

The folder you'll work in

Each knowledge base is its own folder under KNOWLEDGE/. The template ships with eg `_productivity_kb/` as a worked example you can keep, customise, or replace. Your own KBs follow the convention `[topic]_kb` — lowercase, snake_case, ending in `_kb`. For example: `ai_research_kb`, `business_systems_kb`, `theatre_kb`.

Your job is small

You only do three things. Add sources to RAW/. Ask Claude to compile. Ask the KB questions. Everything else — writing wiki articles, cross-linking topics, maintaining INDEX and QUESTIONS files, auditing for drift, drafting new articles — is what the librarian does for you.

Two ways to add sources

Low-token: drop and forget

Drop files straight into `[topic]_kb/RAW/` in Finder. Articles, PDFs, transcripts, screenshots — anything. The librarian processes them next time you ask it to compile. Cheapest path — no chat tokens spent on ingestion. Best for batches.

High-token: guided ingest

Paste sources into chat or share links and ask Claude to ingest them. The librarian adds the required metadata, can ask you context questions if the source is rich, and confirms as it goes. Best for sources where the framing matters — a podcast where the key idea is in passing, a paper with multiple threads worth separating.

Mix and match. Most weeks you'll drop a batch into RAW and ask Claude to compile. Some sources deserve the guided treatment. Both end up in the same place.

Every question becomes a report

When you ask the KB anything, the answer goes into Outputs/ as a dated report, with citations back to the wiki articles it drew from. A short summary lands in chat with a clickable link to the full report. Over time those reports accumulate into a real archive of your thinking — not chat scrollbar that vanishes when you close the window.

Monthly health checks, on autopilot

A skill called `knowledge-base-health-check-skill` ships in RESOURCES/MY SKILLS/. It runs on demand (just say "run a health check") and on the 1st of every month via a scheduled task. Each run audits one KB, auto-fixes routine drift, auto-drafts up to three new wiki articles where there's enough evidence, and flags only the judgement calls that need your input.



GETTING STARTED

Three moves to make today

The fastest way to learn this system is to use it on something you actually care about. Pick a topic you've been collecting on — anything from product strategy to your favourite hobby — and run through these three moves.

1

Spin up a knowledge base

Say to Claude: *"Spin up a knowledge base for [topic]."* Claude will create the folder, copy the template, ask you to confirm the focus areas, and offer to set up the monthly health check on the spot. Use the bundled `eg_productivity_kb/` as a worked example if you want to see one in action first.

2

Drop sources into RAW

Two ways. **Low-token:** drop files straight into `[topic]_kb/RAW/` in Finder — articles, PDFs, transcripts, anything. **High-token:** paste sources into chat or share links and ask Claude to ingest them with metadata. Either works. The verbatim rule is the only thing to remember: sources land word-for-word, never summarised. Then say *"compile"* — Claude reads what's new and writes the wiki.

3

Ask questions, get reports

Once the wiki has shape, ask it anything: *"What does my KB say about X?"* or *"Where do my sources disagree on Y?"* Claude answers from the wiki first, then RAW, with citations. Every answer becomes a dated report in Outputs/ you can come back to.

COMPANION VIDEO

Watch the walkthrough on YouTube

I've recorded a full walkthrough of the knowledge base system — what it does, how I use it, and the tradeoffs to watch out for. If you're new to this pattern, watch the video before your first build.

→ youtu.be/ib74sLgjlBM

REFERENCE

Original Karpathy post: x.com/karpathy/status/2039805659525644595

Top-level rules: `KNOWLEDGE/CLAUDE.md`

Per-KB template: `KNOWLEDGE/_KB_CLAUDE_TEMPLATE.md` — copied into each new KB and customised with the topic and focus areas.